



Last Updated on May 28, 2026

Yichi Zhang

Ph.D Candidate @ Tsinghua University

✉ zyc22@mails.tsinghua.edu.cn, tibo_ricky@outlook.com  [zycheiheiei.github.io](https://github.com/zycheiheiei)

📍 FIT Building 1-508, Tsinghua University, Beijing, China, 100084

EDUCATION

Tsinghua University, Department of Computer Science and Technology August 2022 – June 2027
Beijing, China

Ph.D candidate advised by Prof. Jun Zhu

Research interest in **(Multimodal) LLM Alignment and Evaluation**

Currently researching **reliable agentic system in real-world applications**

Tsinghua University, Department of Computer Science and Technology August 2018 – July 2022
Beijing, China

Bachelor of Engineering, GPA: 3.90/4.00, Ranking: 7/235

Secondary Bachelor of Science in Psychology

NCEE: 709/750, **8th top scorer** of science in Beijing (~35k students)

PUBLICATIONS

(* indicates equal contribution)

PUBLISHED IN CONFERENCES

GLEAN: Guideline-Grounded Evidence Accumulation for High-Stakes Agent Verification

Yichi Zhang, Nabeel Seedat, Yinpeng Dong, Peng Cui, Jun Zhu, Mihaela van de Schaar

International Conference on Machine Learning (ICML), 2026

Best Paper Award at the workshop of Agentic AI in the Wild in ICLR 2026

Towards Safe Reasoning in Large Reasoning Models via Corrective Intervention

Yichi Zhang, Yue Ding, Jingwen Yang, Tianwei Luo, Dongbai Li, Ranjie Duan, Qiang Liu, Hang Su, Yinpeng Dong, Jun Zhu

International Conference on Learning Representations (ICLR), 2026

STAIR: Improving Safety Alignment with Introspective Reasoning (**Oral, ~top 0.9%**)

Yichi Zhang*, Siyuan Zhang*, Yao Huang, Zeyu Xia, Zhengwei Fang, Xiao Yang, Ranjie Duan, Dong Yan, Yinpeng Dong, Jun Zhu

International Conference on Machine Learning (ICML), 2025

RealSafe-R1: Safety-Aligned DeepSeek-R1 without Compromising Reasoning Capability

Yichi Zhang, Zihao Zeng, Dongbai Li, Yao Huang, Zhijie Deng, Yinpeng Dong

R2-FM Workshop at International Conference on Machine Learning (ICML), 2025

MULTITRUST: A Comprehensive Benchmark Towards Trustworthy Multimodal Large Language Models

Yichi Zhang*, Yao Huang*, Yitong Sun, Chang Liu, Zhe Zhao, Zhengwei Fang, Yifan Wang, Huanran Chen, Xiao Yang, Xingxing Wei, Hang Su, Yinpeng Dong, Jun Zhu

Advances in Neural Information Processing Systems (NeurIPS), 2024

Exploring the Transferability of Visual Prompting for Multimodal Large Language Models (**Highlight, ~top 2.8%**)

Yichi Zhang, Yinpeng Dong, Siyuan Zhang, Tianzan Min, Hang Su, Jun Zhu

IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2024

Understanding the Robustness of 3D Object Detection With Bird’s-Eye-View Representations in Autonomous Driving

Zijian Zhu*, **Yichi Zhang***, Hai Chen, Yinpeng Dong, Shu Zhao, Wenbo Ding, Jiachen Zhong, Shibao Zheng

IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2023

MESA: Improving MoE Safety Alignment via Decentralized Expertise

Yitong Sun, Yao Huang, Teng Li, Ranjie Duan, **Yichi Zhang**, Xingjun Ma, Hui Xue, Xingxing Wei

International Conference on Machine Learning (ICML), 2026

Unveiling the Basin-Like Loss Landscape in Large Language Models

Huanran Chen, Yinpeng Dong, Zeming Wei, Yao Huang, **Yichi Zhang**, Hang Su, Jun Zhu

International Conference on Learning Representations (ICLR), 2026

DeceptionBench: A Comprehensive Benchmark for AI Deception Behaviors in Real-world Scenarios

Yao Huang, Yitong Sun, **Yichi Zhang**, Ruochen Zhang, Yinpeng Dong, Xingxing Wei

Advances in Neural Information Processing Systems (NeurIPS), 2025

Mitigating Overthinking in Large Reasoning Models via Manifold Steering

Yao Huang, Huanran Chen, Shouwei Ruan, **Yichi Zhang**, Xingxing Wei, Yinpeng Dong

Advances in Neural Information Processing Systems (NeurIPS), 2025

Exploring the Generalizability of Factual Hallucination Mitigation via Enhancing Precise Knowledge Utilization

Siyuan Zhang, **Yichi Zhang**, Yinpeng Dong, Hang Su

The 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP), Findings, 2025

Breaking the Ceiling: Exploring the Potential of Jailbreak Attacks through Expanding Strategy Space

Yao Huang, Yitong Sun, Shouwei Ruan, **Yichi Zhang**, Yinpeng Dong, Xingxing Wei

Annual Meeting of the Association for Computational Linguistics (ACL), Findings, 2025

PINNacle: A Comprehensive Benchmark of Physics-Informed Neural Networks for Solving PDEs

Zhongkai Hao, Jiachen Yao, Chang Su, Hang Su, Ziao Wang, Fanzhi Lu, Zeyu Xia, **Yichi Zhang**, Songming Liu, Lu Lu, Jun Zhu

Advances in Neural Information Processing Systems (NeurIPS), 2024

Rethinking Model Ensemble in Transfer-based Adversarial Attacks

Huanran Chen, **Yichi Zhang**, Yinpeng Dong, Jun Zhu

International Conference on Learning Representations (ICLR), 2024

PUBLISHED IN JOURNALS

To make yourself invisible with Adversarial Semantic Contours

Yichi Zhang, Zijian Zhu, Hang Su, Jun Zhu, Shibao Zheng, Yuan He, Hui Xue

Computer Vision and Image Understanding (CVIU), 2023

RELEASED AS PREPRINTS

Guideline-Grounded Evidence Accumulation for High-Stakes Agent Verification

Yichi Zhang, Nabeel Seedat, Yinpeng Dong, Peng Cui, Jun Zhu, Mihaela van de Schaar

ArXiv, 2026

Unveiling Trust in Multimodal Large Language Models: Evaluation, Analysis, and Mitigation

Yichi Zhang, Yao Huang, Yifan Wang, Yitong Sun, Chang Liu, Zhe Zhao, Zhengwei Fang, Huanran Chen, Xiao Yang, Xingxing Wei, Hang Su, Yinpeng Dong, Jun Zhu

ArXiv (Under Review for TPAMI), 2025

Physics-informed machine learning: A survey on problems, methods and applications
Zhongkai Hao, Songming Liu, **Yichi Zhang**, Chengyang Ying, Yao Feng, Hang Su, Jun Zhu
ArXiv, 2022

EXPERIENCE

Hunyuan, Tencent <i>Qingyun Research Intern</i> Coding Agent for Cybersecurity	April 2026 – Present
University of Cambridge <i>Visiting PhD Student</i> Agent Verification in Healthcare	October 2025 – April 2026
Qwen Team, Alibaba <i>Research Intern</i> Agentic System for DeepResearch	May 2025 – August 2025
University of Sydney <i>Visiting PhD Student</i> Trustworthy Machine Learning	March 2025 – May 2025
RealAI <i>Research Intern</i> Safety and robustness of deep learning models in wide applications	January 2023 – December 2024
Tencent <i>Research Intern</i> Advertising models and re-ranking models in the recommendation system of Tencent Video Platform	July 2021 – September 2021

COMPETITIONS

The 1st place in the Adversarial Robustness track of 2022 International Algorithm Case Competition	February 2023
The 2nd place in the CVPR 2021 Security AI Challenger Unrestricted Adversarial Attacks on ImageNet	June 2021
The 8th place in the CIKM 2020 Adversarial Challenge on Object Detection	September 2020

SELECTED AWARDS

Tencent PhD Research Incentive Program (“ <i>Hunyuan Scholar</i> ”, 23 recipients nationwide)	July 2025
Tsinghua Outstanding Graduates (top 2%)	June 2022
Beijing Outstanding Graduates (top 5%)	June 2022
Beijing Merit Students (Only one student in the department each year)	March 2022
Tsinghua Overall Excellence Scholarships	December 2019,2020,2021,2024

SERVICE

Organizer

ICML2024 Workshop on Trustworthy Multi-modal Foundation Models and AI Agents (TiFA)
CVPR2025 Workshop on Test-time Scaling for Computer Vision (ViSCALE)
ICCV2025 Workshop on Safe and Trustworthy Multimodal AI Systems (SaFeMM-AI)
CVPR2026 Second Workshop on Test-time Scaling for Computer Vision (ViSCALE)

Reviewer

ICML, ICLR, NeurIPS, CVPR, ICCV, ECCV, ACL, EMNLP, TPAMI

Teaching

TA in Machine Learning, instructed by Prof. Jun Zhu and Prof. Jie Tang, 2023 Autumn